

```
In [1]: import pyspark
import findspark
from pyspark.sql import SparkSession

from pyspark.ml.feature import VectorAssembler
from pyspark.ml.classification import LogisticRegression

import pandas as pd

from pyspark.ml.evaluation import BinaryClassificationEvaluator
```

```
In [2]: spark = SparkSession.builder.master('local[*]').getOrCreate()
```

```
23/04/20 01:22:39 WARN Utils: Your hostname, MacBook-Air-Evgenij.local resolves to a loopback address: 127.0.0.1; using 192.168.0.102 instead (on interface en0)
23/04/20 01:22:39 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
23/04/20 01:22:39 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

```
In [3]: spark
```

```
Out[3]: SparkSession - in-memory
```

SparkContext

[Spark UI](#)

Version	v3.4.0
Master	local[*]
AppName	pyspark-shell

```
In [4]: df = pd.read_csv('./santander-customer-transaction-prediction/train.csv')
```

```
In [5]: df.head()
```

```
Out[5]:
```

	ID_code	target	var_0	var_1	var_2	var_3	var_4	var_5	var_6	va
0	train_0	0	8.9255	-6.7863	11.9081	5.0930	11.4607	-9.2834	5.1187	18.62
1	train_1	0	11.5006	-4.1473	13.8588	5.3890	12.3622	7.0433	5.6208	16.53
2	train_2	0	8.6093	-2.7457	12.0805	7.8928	10.5825	-9.0837	6.9427	14.6
3	train_3	0	11.0604	-2.1518	8.9522	7.1957	12.5846	-1.8361	5.8428	14.92
4	train_4	0	9.8369	-1.4834	12.8746	6.6375	12.2772	2.4486	5.9405	19.2

5 rows × 202 columns

```
In [6]: df.dtypes
```

```
Out[6]: ID_code      object
target      int64
var_0       float64
var_1       float64
var_2       float64
...
var_195     float64
var_196     float64
var_197     float64
var_198     float64
var_199     float64
Length: 202, dtype: object
```

```
In [7]: input_columns = df.columns.tolist()[2:]
```

```
In [8]: dataset = spark.read.csv('./santander-customer-transaction-prediction/train.csv')
```

```
In [9]: input_columns = dataset.columns[2:]
```

```
In [10]: assembler = VectorAssembler(inputCols=input_columns, outputCol='Attributes')
output = assembler.transform(dataset)
finalized_data = output.select('Attributes', 'target')
finalized_data.show()
```

23/04/20 01:22:46 WARN package: Truncated the string representation of a plan since it was too large. This behavior can be adjusted by setting 'spark.sql.debug.maxToStringFields'.

Attributes	target
[8.9255,-6.7863,1...	0
[11.5006,-4.1473,...	0
[8.6093,-2.7457,1...	0
[11.0604,-2.1518,...	0
[9.8369,-1.4834,1...	0
[11.4763,-2.3182,...	0
[11.8091,-0.0832,...	0
[13.558,-7.9881,1...	0
[16.1071,2.4426,1...	0
[12.5088,1.9743,8...	0
[5.0702,-0.5447,9...	0
[12.7188,-7.975,1...	0
[8.7671,-4.6154,9...	0
[16.3699,1.5934,1...	1
[13.808,5.0514,17...	0
[3.9416,2.6562,13...	0
[5.0615,0.2689,15...	0
[8.4199,-1.8128,8...	0
[4.875,1.2646,11....	0
[4.409,-0.7863,15...	0

only showing top 20 rows

```
In [11]: train_data, test_data = finalized_data.randomSplit([0.8, 0.2])

regressor = LogisticRegression(featuresCol='Attributes', labelCol='target')

regressor = regressor.fit(train_data)

pred = regressor.evaluate(test_data)

pred.predictions.show()
```

```
23/04/20 01:22:51 WARN InstanceBuilder: Failed to load implementation from:de
v.ludovic.netlib.blas.JNIBLAS
23/04/20 01:22:51 WARN InstanceBuilder: Failed to load implementation from:de
v.ludovic.netlib.blas.VectorBLAS
23/04/20 01:22:52 WARN GarbageCollectionMetrics: To enable non-built-in garba
ge collector(s) List(G1 Concurrent GC), users should configure it(them) to sp
ark.eventLog.gcMetrics.youngGenerationGarbageCollectors or spark.eventLog.gcM
etrics.oldGenerationGarbageCollectors
```

Attributes	target	rawPrediction	probability	predic tion
------------	--------	---------------	-------------	----------------

```

----+
| [1.6044,-0.6302,1... | 0 | [2.81577733270635... | [0.94352247066797... |
0.0 |
| [1.922,-3.31,8.37... | 0 | [2.97425341712561... | [0.95139733495655... |
0.0 |
| [2.5041,-4.9081,8... | 0 | [5.21681693215687... | [0.99460469873667... |
0.0 |
| [2.6501,-2.6365,1... | 0 | [3.24585880260540... | [0.96252401910992... |
0.0 |
| [2.8253,4.2618,15... | 0 | [5.54884902547382... | [0.99612315424485... |
0.0 |
| [3.6355,4.1567,14... | 1 | [1.94706342408893... | [0.87512608491385... |
0.0 |
| [3.7917,-1.0115,8... | 0 | [4.41978174446677... | [0.98810630363823... |
0.0 |
| [3.8296,-7.5727,8... | 0 | [3.69750919841062... | [0.97581426332228... |
0.0 |
| [3.8649,2.1927,10... | 0 | [2.66281427774470... | [0.93479641269697... |
0.0 |
| [3.9177,1.2849,5... | 0 | [6.25206108581090... | [0.99807722478470... |
0.0 |
| [3.9977,2.2683,8... | 0 | [3.19744450777572... | [0.96073799617652... |
0.0 |
| [4.0771,-7.9056,7... | 0 | [5.50612310511931... | [0.99595460751027... |
0.0 |
| [4.1896,-2.8807,1... | 1 | [0.61723846028127... | [0.64959021922551... |
0.0 |
| [4.203,-7.5025,14... | 0 | [4.14977376339386... | [0.98447678639762... |
0.0 |
| [4.2341,-8.5367,1... | 0 | [5.33895359777056... | [0.99522204668995... |
0.0 |
| [4.2485,1.1863,6... | 0 | [4.03553113542261... | [0.98263073503312... |
0.0 |
| [4.3542,2.359,9.6... | 0 | [6.02678492172501... | [0.99759256943255... |
0.0 |
| [4.3836,-1.4746,9... | 0 | [2.28617095563942... | [0.90772522910623... |
0.0 |
| [4.3854,-3.0492,7... | 0 | [1.47928493562790... | [0.81446455016284... |
0.0 |
| [4.409,-0.7863,15... | 0 | [3.93850733776943... | [0.98089485005077... |
0.0 |
+-----+-----+-----+-----+

```

----+
only showing top 20 rows

```
In [12]: coef = regressor.coefficients
```

```
intr = regressor.intercept
```

```
print('coef: %a' % coef)
```

```
print('intr: %a' % intr)
```

```
coef: DenseVector([0.0531, 0.0386, 0.0591, 0.0221, 0.0224, 0.0147, 0.2342, -0.0038, 0.0217, -0.084, 0.0019, 0.0161, -1.1831, -0.0341, -0.0037, 0.167, 0.0077, -0.0016, 0.0196, 0.0072, -0.011, -0.0235, 0.0629, -0.1602, 0.0346, 0.1419, 0.0388, -0.0081, -0.1076, 0.0057, -0.0006, -0.0439, 0.0314, -0.0396, -0.3118, 0.0262, -0.0395, 0.008, -0.0032, -0.0026, 0.0187, -0.0011, -0.0406, -0.3514, -0.0319, -0.0022, 0.002, 0.0031, 0.0093, 0.0109, -0.0538, 0.0048, 0.0153, 0.2585, -0.0067, 0.0129, -0.0269, -0.0896, -0.0188, -0.0352, 0.0098, 0.0038, 0.0228, -0.0182, -0.0302, 0.01, 0.0784, 0.0206, -5.3815, 0.0062, 0.0056, 0.3775, -0.0163, -0.0008, 0.0059, -0.0266, -0.0236, -0.0115, 0.0718, 0.0341, -0.0283, -0.1025, 0.0099, -0.0064, 0.0099, -0.0256, -0.0193, -0.0147, -0.0225, 0.0341, 0.0061, 0.8587, -0.0326, -0.2586, 0.0663, 0.2246, 0.0026, 0.0058, -0.0157, 0.1014, -0.0009, -0.0107, -0.0068, 0.113, -0.039, 0.1175, 0.065, -0.0161, -0.8179, -0.035, 0.0521, 0.0885, 0.0538, -0.0201, -0.109, -0.0605, -0.0562, 0.0009, 0.015, 0.0238, -0.004, -0.0736, -0.03, -0.0208, -0.0065, 0.245, 0.029, -0.0402, 0.0269, -0.0061, 0.1545, -0.2222, -0.0642, 0.422, 0.0073, 0.0116, -0.0027, 0.0089, 0.0093, -0.0278, 0.0134, -0.0128, -0.0164, -0.0207, 0.1115, 0.0295, -0.0843, 0.0185, -0.8814, -0.0132, -0.0298, 0.026, -0.0128, -0.0105, -0.0268, 0.0219, -0.0733, 0.0127, -0.004, 0.0157, -0.0007, 0.1049, 0.0831, 0.0196, 0.0249, -0.0376, -0.5367, 0.011, 0.013, -0.4508, 0.0379, 0.0116, -0.0147, 0.0286, -0.0261, 0.0355, 0.0054, -0.0408, -0.0043, 0.0594, 0.0194, 0.0416, -0.004, 0.0009, 0.017, 0.0011, -0.0309, 0.004, -0.0308, 0.0435, 0.0397, 0.0429, -0.0814, -0.016, -0.0215, 0.0638, 0.0154, -0.1337, -0.0572, 0.0097])
intr: 60.6480039546467
```

```
In [13]: eval = BinaryClassificationEvaluator(rawPredictionCol='rawPrediction', labelCol='label')
```

```
In [14]: roc = eval.evaluate(pred.predictions)
print('ROC: %.3f' % roc)
```

```
ROC: 0.853
```